



Summarizer

User Guide

Record, transcribe and summarize meetings with AI
macOS / Whisper / Gemini / GPT-5 / Ollama

v1.18.4

Installation

One-time setup — takes about a minute.

1 Download Summarizer.dmg

Go to github.com/fvoin/summarizer/releases and download the latest DMG file.

2 Open the DMG

Double-click the downloaded file. A window opens showing Summarizer and an Applications folder.

3 Drag to Applications

Drag the Summarizer icon into the Applications folder shortcut.

4 Launch and allow

Open Summarizer from Applications (Launchpad or Finder).

macOS will block it — right-click the app → Open → click Open again.


Or: System Settings → Privacy & Security → scroll down → Open Anyway.

5 Set up an AI model

Click the gear icon → Models tab.

Cloud: enter your API key (free Gemini key at aistudio.google.com/apikey)

Local: click Download on any Ollama model — no API key, fully offline.



First launch only (bypass macOS quarantine):

Option A: Right-click the app → Open → click Open in the dialog

Option B: System Settings → Privacy & Security → Open Anyway

This is required once for any app not from the App Store.

What is Summarizer?

Summarizer is a macOS application that automatically records, transcribes, and summarizes meetings and audio recordings.

It uses **Whisper** for local speech recognition (audio never leaves your machine) and LLM models (Gemini, Claude, OpenAI) to generate structured summaries.

Key Features

Audio recording — record from your microphone with automatic silence detection and auto-stop

Local transcription — Whisper models (tiny through large-v3) run directly on your Mac, no cloud needed

AI summarization — structured output: Overview, Key Decisions, Action Items, Discussion Points, Risks, and a Meeting Score with cost estimate

Fully offline mode — use local LLM models (Llama, Gemma, Qwen) via Ollama for 100% offline operation — nothing ever leaves your machine

Instruction profiles — create multiple prompt profiles and switch between them for different meeting types (standup, review, 1-on-1...)

Meeting series — accumulate meeting history for more accurate summaries across recurring meetings

Meeting history — SQLite database stores all transcripts, summaries, and contexts with browsable history

Meeting series chat — ask questions about your meeting history using any configured LLM

Workflow



The entire process takes from a few seconds to a couple of minutes

Privacy

Audio is processed **LOCALLY** by Whisper — it never leaves your machine.


With cloud models, only the text transcript is sent for summarization.

With local models (Ollama), everything stays on your Mac — fully offline.

No data is collected. No accounts required. You own all your data.

Main Window

Everything is accessible from a single window.



Recording Audio

1 Press Start Recording

The button turns red and shows a timer. The app records audio from the selected input device (configured in Settings).



2 Automatic Stop

Recording stops automatically after a period of silence (configurable in Settings, default 180 seconds / 3 minutes). You can also stop manually by pressing the red button.

3 Transcription

After stopping, Whisper automatically recognizes the speech. The model runs locally — audio never leaves your machine. Progress is shown with a color-coded status indicator.

4 Summary & Result

The LLM generates a structured summary. Use Copy Summary to copy with formatting that works in Slack. The transcript is saved to a file you can open.

Tips

- Whisper runs locally — audio is never sent to the cloud
- For better transcription quality, use the medium or large-v3 model
- Saving audio files can be enabled in Settings (off by default)

Menu Bar & Recording Agent

Menu Bar Icon

Summarizer places an icon in the macOS menu bar. Clicking it reveals a compact menu with **Start/Stop Recording**, **Show Summarizer**, **Settings**, and **Quit**.

Hide from Dock

When minimized to the menu bar tray, the app hides from the Dock. Click the menu bar icon and choose **Show Summarizer** to bring it back.


Single Instance

Launching Summarizer again while it is already running does not open a second copy. Instead, the existing window is brought to the front.

Recording Agent

The Recording Agent automatically starts recording when an upcoming meeting is detected from a web calendar backend. Configure the backend URL and behavior in **Settings** → **Recording Agent**.

When a meeting is about to start, the agent activates recording and stops when the meeting ends. The transcript and summary are generated automatically, just like a manual recording session.



Menu bar dropdown

Meeting History & Series

Meeting History

Click the **clock icon** in the top bar to open the Meeting History dialog. It shows a searchable, scrollable list of all past meetings with their date, series, transcript, and summary. All data is stored in a local SQLite database.

Meeting Series Editor

Click the **edit button** (pencil) next to the series dropdown to open the series editor. Here you can rename series, update general context, review accumulated history, and delete series you no longer need.

Series Chat

Click the **chat icon** (chat bubble) next to the series dropdown to open a chat window scoped to that meeting series. Ask questions about past meetings, action items, decisions, or trends — the LLM has access to the full series history stored in the database.

SQLite Storage

All transcripts, summaries, meeting contexts, and series data are stored in a local SQLite database at **~/.summarizer/summarizer.db**. This replaces the older plain-text context files and provides faster search, reliable history, and structured queries for the series chat feature.

~/.summarizer/

summarizer.db — SQLite database (meetings, series, contexts)

recordings/ — audio files (if Save Audio is enabled)

Settings — Models

Choose between cloud and local AI models for summarization.

Settings

General

Models

Instructions

Recording Agent

☐ Cloud

☒ Gemini 3 Flash Preview

☐ Gemini 2.5 Pro

☐ GPT-5 mini

☐ GPT-5.4

☐ Custom:

API Key:

Base URL:

☐ Local (Ollama)

☐ GLM-4 9B

— Good (5.5 GB)

Download

☐ Gemma 3 12B QAT

— Better (8.9 GB)

Download

☐ Qwen 3 30B

— Great (19 GB)

Download

☒ GPT-OSS 20B

— Best (12 GB)

ReadyTestDelete

Whisper Model (speech recognition)

☐ tiny

— Basic (75 MB)

Download

☒ base

— Good (145 MB)

Ready

☐ small

— Better (465 MB)

Download

☐ medium

— Great (1.5 GB)

Download

☐ large-v3

— Best (3.1 GB)

Download

Save

Cancel

Cloud presets

Select a model or enter a custom model name

API credentials

Key for the selected cloud provider

Local models (Ollama)

Run AI on your Mac — 100% offline, no API key. Download, Test, or Delete.

Whisper models

Speech-to-text (offline). Larger = more accurate but slower.

Settings — Instructions

Customize the AI prompt. Create profiles for different meeting types.

Settings

General

Models

Instructions

Recording Agent

Profile:

Default

▼

New

Delete

You are a professional meeting analyst. Produce a structured, actionable summary of the transcript below.

Output exactly these sections in order:

☐

Overview

• One sentence: meeting purpose and main outcome.

☐

Key Decisions

• Each confirmed decision, stated as a fact.

• If none — omit this section entirely.

☐

Action Items

• Format: ***Owner*** — task — **_deadline if mentioned_**

• If owner is unclear, write **_unassigned_**.

☐

Key Discussion Points

• Important topics discussed, options considered.

☐

Risks & Open Questions

• Unresolved issues, blockers, follow-ups.

☐

Meeting Score

• ***Decision output***: X/10

• ***Time efficiency***: X/10

• ***Actionability***: X/10

• ***Final***: avg of 3 — X/10

• ***Could this have been an email?*** Yes/No

• ***Cost estimate***: [duration]h × [N] × 50 EUR = ~X EUR

FORMATTING: • only, ***bold***, *_italic_*, no # markdown

Save

Cancel

Profiles

Create separate prompts for standups, reviews, 1-on-1s. Switch with the dropdown.

Default prompt

Structured meeting summary with sections: Overview, Decisions, Actions, Score. Edit freely or replace.

Formatting rules

Controls how the LLM formats output — bullets, bold, italic, no markdown.

Settings — General

Audio, recording, and storage preferences.

Settings

General

Models

Instructions

Recording Agent

Context Limit:

5000 chars

Max characters loaded from context file.
Older content is trimmed automatically.

Silence Timeout:

180 sec

Auto-stop after this much silence (3 min).
Adaptive calibration measures mic noise first.

Input Device:

Default

Choose your microphone or loopback device.
Default uses the system input.

Save Audio:

off

When on, saves WAV recordings to disk.
Off by default — recordings are temporary.

Sound:

on

Play a sound when the summary is ready.

Recordings Dir:

~/summarizer/recordings

Where transcripts and context files are stored.

Version:

v1.18.4

Check for Updates

Open Log File

~/summarizer/summarizer.log

Save

Cancel

Meeting Series

What is a Meeting Series?

A meeting series groups recurring meetings (standups, 1-on-1s, reviews) so Summarizer can track progress and action items across sessions. All data is stored in a local SQLite database.

Creating a Series

Click **+** next to the Meeting series dropdown and enter a name. Use **edit (pencil)** to open the series editor, **chat (bubble)** to chat about the series, and **×** to delete it.

Persistent Context

Each series has a persistent context — participants, goals, key terms. Edit it via the **edit (pencil)** series editor. This is always included in the LLM prompt for accurate summaries.

This Meeting Context

Per-session details — today's agenda, attendees, specific topics. Always included in the prompt. Saved alongside the transcript and summary in the database.

History & Budget

Each meeting's summary (with its meeting context) is stored in the database. The Context Limit (default 5000 chars) controls how much history is loaded into the prompt — persistent context and current meeting context are always included in full, remaining budget is filled with recent summaries.

Series Chat

Click **chat (bubble)** to chat with an AI about the series. The model sees the persistent context and all recent meeting summaries. Ask questions like 'what did we decide about X?' or 'what are the open action items?'

Data Storage

SQLite database: `~/summarizer/summarizer.db`

contexts table — series name, persistent context

meetings table — transcript, summary, meeting context,
duration, date, linked to series

Existing `_context.txt` files are auto-migrated on first run.

Browse all data via the History dialog (📄 icon in top bar).

FAQ & Tips

Where do I get an API key?

Gemini: aistudio.google.com - Create API Key. Claude: console.anthropic.com - API Keys. OpenAI: platform.openai.com - API Keys.

Formatting doesn't work when pasting to Slack?

Use the **Copy Summary** button — it copies text with HTML markup, so bold and italic render correctly. Regular Cmd+C copies plain text only.

Which Whisper model should I choose?

base (145 MB) — fast, good enough for clear speech. **medium** (1.5 GB) — best balance of speed and accuracy. **large-v3** (3.1 GB) — maximum accuracy, but slower.

Where is data stored?

Database: `~/summarizer/summarizer.db` (meetings, contexts, transcripts). Config: `~/summarizer/config.json`. Whisper models: `~/summarizer/models/`. App log: `~/summarizer/summarizer.log`.

Is it safe?

Audio is always processed **locally** by Whisper — it never leaves your Mac. With cloud models, only the text transcript is sent to the LLM. With **local models (Ollama)**, everything stays on your Mac — fully offline.

How do I use the menu bar mode?

Enable **Menu bar icon** in Settings > Recording Agent. Close the window — the app stays in the menu bar. Click the tray icon or launch the app again to show the window.

Can I use a custom LLM endpoint?

Yes — set the **Base URL** field in Settings to point to any OpenAI-compatible API endpoint.